

COT 6405 Introduction to Theory of Algorithms

Topic 10. Linear Time Sorting

How fast can we sort?

- The sorting algorithms we learned so far
 - Insertion Sort, Merge Sort, Heap Sort, and Quicksort
- How fast are they?
 - Insertion sort $O(n^2)$
 - Merge Sort $O(n \lg n)$
 - Heap Sort $O(n \lg n)$
 - Quicksort $O(n \lg n)$

Common property


- Use only comparisons between elements to gain order information about an input sequence
- Comparison sort
 - Given two elements a_i and a_j , we perform one of the following tests to determine their relative order
 - $a_i < a_j$, $a_i \leq a_j$, $a_i = a_j$, $a_i \geq a_j$, $a_i > a_j$

Decision trees

- We can view comparison sorts abstractly in terms of decision trees
 - A decision tree is a binary tree that represents the comparisons between elements
 - Each node on the tree is a comparison of $i:j$, i.e., a_i v.s. a_j

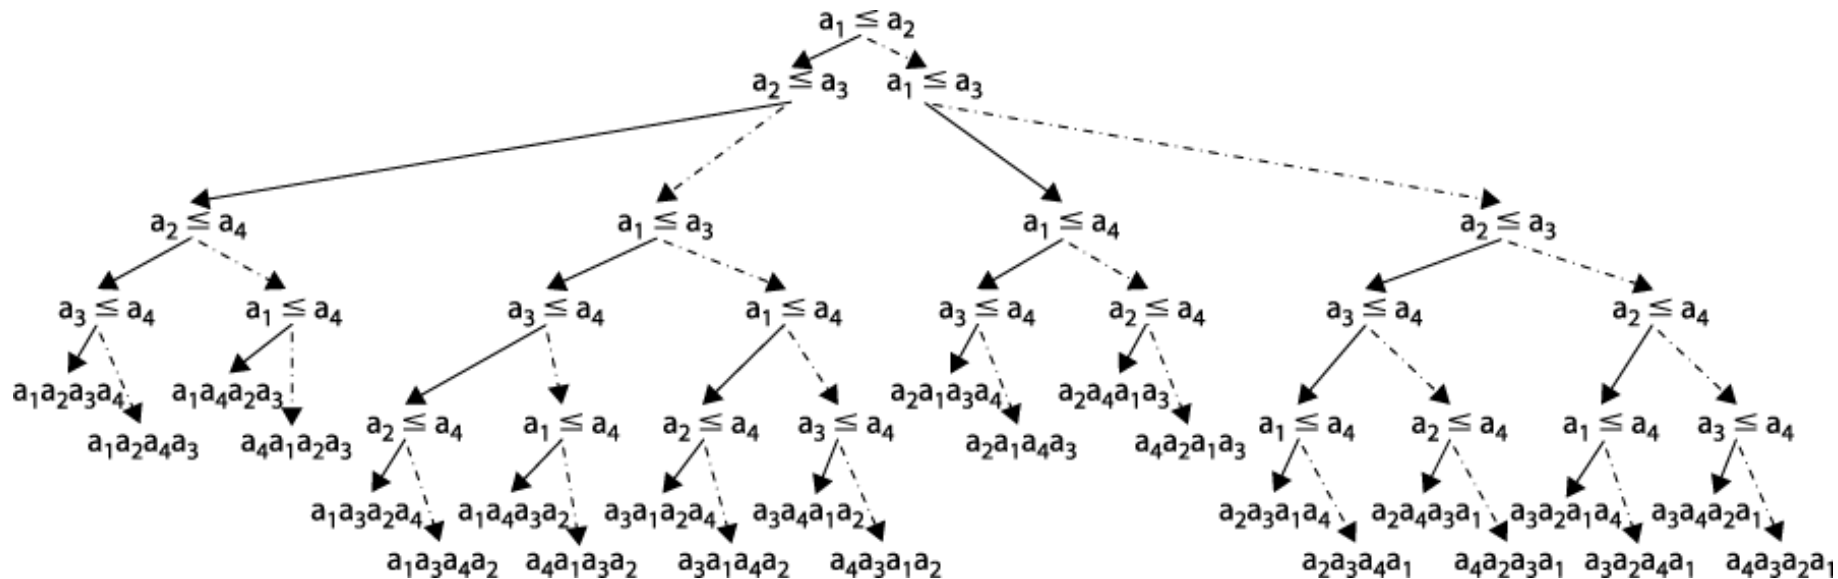
Constructing the decision tree

- Given an input sequence $\{a_1, a_2, a_3\}$



Decision tree for an input set of four elements

Given an input sequence $\{a_1, a_2, a_3, a_4\}$



Decision trees (cont'd)

- What do the leaves represent?
 - The leaf node in the tree indicates the sorted ordering
- How many leaves must be there for an input of size n
 - Each of the $n!$ permutations on n elements must appear as one of the leaves of the decision tree

Lemma

- Any binary tree of height h has $\leq 2^h$ leaves
- In other words:
 - i = number of leaves
 - h = height
 - Then, $i \leq 2^h$
- How to prove this?

Theorem 8.1

- Any comparison sort algorithm requires $\Omega(n \lg n)$ comparisons.
- How to prove?
 - By proving that the height of the decision tree is $\Omega(n \lg n)$
 - What's the # of leaves of a decision tree? $l = ?$
 - A decision tree is a binary tree. What's the maximum # of leaves of a general binary tree?
 $l_{\max} = ?$

Proof

- $I = n!$ and $I_{\max} = 2^h$
- Clearly, the # of leaves of a decision tree is less than or equal to the maximum # of leaves in a general binary tree
- So we have: $n! \leq 2^h$
- Taking logarithms: $\lg(n!) \leq h$

Proof (cont'd)

- Stirling's approximation tells us:

$$n! > \left(\frac{n}{e}\right)^n$$

- Thus, $h \geq \lg(n!)$

$$\begin{aligned} h &\geq \lg \left(\frac{n}{e}\right)^n \\ &= n \lg n - n \lg e \\ &= \Omega(n \lg n) \end{aligned}$$

Sorting in linear time

- Counting sort
 - No direct comparisons between elements!
 - Depends on assumption about the numbers being sorted
 - We assume numbers are in the range $[0..k]$
 - The algorithm is NOT “in place”
 - Input: $A[1..n]$, where $A[j] \in \{0, 2, 3, \dots, k\}$
 - Output: $B[1..n]$, sorted
 - Auxiliary counter storage: Array $C[0..k]$
 - notice: $A[]$, $B[]$, and $C[] \rightarrow$ not sorting in place

Counting sort

```
1  CountingSort(A, B, k)
2      for i= 0 to k    // counter initialization
3          C[i]= 0;
4      for j= 1 to A.length
5          C[A[j]] += 1;
6      for i= 1 to k    // aggregate counters
7          C[i] = C[i] + C[i-1];
8      for j= A.length downto 1 //move results
9          B[C[A[j]]] = A[j];
10         C[A[j]] -= 1;
```

A counting sort example

Numbers are in the range [0.. 5]

A

2	5	3	0	2	3	0	3
---	---	---	---	---	---	---	---

 $k = 5$

C

0	0	0	0	0	0
---	---	---	---	---	---

```
1 CountingSort(A, B, k)
2   for i= 0 to k    // counter initialization
3       C[i]= 0;
4   for j= 1 to A.length
5       C[A[j]] += 1;
6   for i= 1 to k    // aggregate counters
7       C[i] = C[i] + C[i-1];
8   for j= A.length downto 1 //move results
9       B[C[A[j]]] = A[j];
10      C[A[j]] -= 1;
```

Filling the C array

A	2	5	3	0	2	3	0	3
	0	1	2	3	4	5		
C	2	0	2	3	0	1		

```
1 CountingSort(A, B, k)
2   for i= 0 to k    // counter initialization
3       C[i]= 0;
4   for j= 1 to A.length // counting each number
5       C[A[j]] += 1;
6   for i= 1 to k    // aggregate counters
7       C[i] = C[i] + C[i-1];
8   for j= A.length downto 1 //move results
9       B[C[A[j]]] = A[j];
10      C[A[j]] -= 1;
```

Filling the C array (Cont'd)

	0	1	2	3	4	5
C	2	2	4	7	7	8

```
1 CountingSort(A, B, k)
2   for i= 0 to k    // counter initialization
3       C[i]= 0;
4   for j= 1 to A.length
5       C[A[j]] += 1;
6   for i= 1 to k    // aggregate counters
7       C[i] = C[i] + C[i-1];
8   for j= A.length downto 1 //move results
9       B[C[A[j]]] = A[j];
10      C[A[j]] -= 1;
```


Sorting the numbers

A	2	5	3	0	2	3	0	3
	1	2	3	4	5	6	7	8
B								
	0	1	2	3	4	5		
C	2	2	4	7	7	8		

```
1 CountingSort(A, B, k)
2   for i= 0 to k    // counter initialization
3       C[i]= 0;
4   for j= 1 to A.length
5       C[A[j]] += 1;
6   for i= 1 to k    // aggregate counters
7       C[i] = C[i] + C[i-1];
8   for j= A.length downto 1 //move results
9       B[C[A[j]]] = A[j];
10      C[A[j]] -= 1;
```

Sorting the numbers

A	2	5	3	0	2	3	0	3
	1	2	3	4	5	6	7	8
B		0					3	
	0	1	2	3	4	5		
C	1	2	4	6	7	8		

```
1 CountingSort(A, B, k)
2   for i= 0 to k    // counter initialization
3       C[i]= 0;
4   for j= 1 to A.length
5       C[A[j]] += 1;
6   for i= 1 to k    // aggregate counters
7       C[i] = C[i] + C[i-1];
8   for j= A.length downto 1 //move results
9       B[C[A[j]]] = A[j];
10      C[A[j]] -= 1;
```

Sorting the numbers

A	2	5	3	0	2	3	0	3
	1	2	3	4	5	6	7	8
B		0				3	3	
	0	1	2	3	4	5		
C	1	2	4	5	7	8		


```
1 CountingSort(A, B, k)
2   for i= 0 to k    // counter initialization
3       C[i]= 0;
4   for j= 1 to A.length
5       C[A[j]] += 1;
6   for i= 1 to k    // aggregate counters
7       C[i] = C[i] + C[i-1];
8   for j= A.length downto 1 //move results
9       B[C[A[j]]] = A[j];
10      C[A[j]] -= 1;
```

Counting sort

- Total time: $O(n + k)$
 - Usually, $k = O(n) \rightarrow k < c n$
 - Thus counting sort runs in $O(n)$ time
- But sorting is $\Omega(n \lg n)$! Contradiction?
 - No contradiction--this is not a comparison sort (in fact, there are *no* comparisons at all!)
 - Notice that this algorithm is stable
 - The elements with the same value is in the same order as the original
 - index $i < j$, $a_i = a_j \rightarrow$ new index $i' < j'$

Stable sorting

Counting sort is a stable sort: it preserves the input order among equal elements.



Counting Sort

- Why don't we always use counting sort?
- Because it depends on range k of elements
- Could we use counting sort to sort 32 bit integers? Why or why not?
- Answer: no, k too large ($2^{32} = 4,294,967,296$)
 - We need huge arrays, e.g., $C[4,294,967,296]$?
 - $k \gg n \rightarrow O(n+k) = O(k)$

Radix Sort

- Intuitively, we may sort on the most significant digit (MSD), then the second msd, etc.
- Recursive MSD radix sort:
 - Take the k-th most significant digit (MSD)
 - Sort based on that digit, grouping same digit elements into one bucket
 - In each bucket, start with the next digit and sort recursively
 - Finally, concatenate the buckets in order

An example of a forward recursive MSD radix sort

- Original sequence: 170, 045, 075, 090, 002, 024, 802, 066
- 1st pass- Sorting by most significant digit (100's):
 - Zero bucket: 045, 075, 090, 002, 024, 066
 - One bucket: 170
 - Eight bucket: 802

An example (cont'd)


- 2nd pass- Sorting by next most significant digit (10's), only needed by numbers in zero bucket:
 - 045, 075, 090, 002, 024, 066
 - Zero bucket: 002
 - Twenties bucket: 024
 - Forties bucket: 045
 - Sixties bucket: 066
 - Seventies bucket: 075
 - Nineties bucket: 090

An example (cont'd)

- 3rd pass- Sorting by least significant digit (1's): no need because there are no tens buckets with more than one number.
- 4th pass- The sorted zero hundreds buckets are concatenated and joined in sequence to give 002, 024, 045, 066, 075, 090, 170, 802
 - Zero bucket: 002
 - Twenties bucket: 024
 - Forties bucket: 045
 - Sixties bucket: 066
 - Seventies bucket: 075
 - Nineties bucket: 090
 - Zero bucket: 045, 075, 090, 002, 024, 066
 - One bucket: 170
 - Eight bucket: 802

Most Significant Digit (MSD) Radix Sort

- Problem:
 - lots of intermediate piles of cards to keep track of
 - d digits $\rightarrow T(n) = O(dn)$
 - Needs to use many buckets to store intermediate results, each is a linked list of size up to n
 - MSD sort does not necessarily preserve the original order of duplicate keys
 - Depending on how we sort the bucket



Least significant digit (LSD) Radix Sort


- Key idea: sort the least significant digit first
- Assume we have d -digit numbers in A

```
RadixSort(A, d)
```

```
    for  $i = 1$  to  $d$ 
```

```
        StableSort(A) on digit  $i$ 
```

Example: LSD Radix Sorting



Radix Sort

- Can we prove it works?
- Sketch of an inductive argument (induction on the number of passes)
 - Assume lower-order digits $\{j: j < i\}$ are sorted
 - Show that sorting next digit i leaves array correctly sorted
 - If two digits at position i are different, ordering numbers by that digit is correct (lower-order digits are irrelevant)
 - If they are the same, numbers are already sorted on the lower-order digits. Since we use a stable sort, the numbers stay in the right order

Questions?

- Can we use any sorting algorithms instead of stable sorting in LSD Radix sorting?

Why stable sorting

- 6**5**7 6**5**8 469 595
- If the sorting algorithm is not stable
- First pass: 595 6**5**7 6**5**8 469
- Second pass: 6**5**8 6**5**7 469 595
- Third pass: 469 595 6**5**8 6**5**7

Radix Sort

- What sort will we use to sort on digits?
- Counting sort is obvious choice:
 - Sort n numbers on digits that range from $0..k$
 - Time: $O(n + k)$
- Each pass over n numbers with d digits takes time $O(n+k)$, so the total time $O(dn+dk)$
 - When d is constant and $k = O(n)$, takes $O(n)$ time

How to break words into digits?

- We have n word
- Each word is of b bits
- We break each word into r -bit digits, $d = \lceil b/r \rceil$
- Using counting sort, $k = 2^r - 1$
- E.g., 32-bit word, we break into 8-bit digits
 - $d = \lceil 32/8 \rceil = 4$, $k = 2^8 - 1 = 255$
- $T(n) = \Theta(d*(n+k)) = \Theta(b/r * (n + 2^r))$

How to choose r ?

How to choose r ? Balance b/r and $n + 2^r$. Choosing $r \approx \lg n$ gives us $\Theta\left(\frac{b}{\lg n} (n + n)\right) = \Theta(bn/\lg n)$.

Still in $O(n)$

- If we choose $r < \lg n$, then $b/r > b/\lg n$, and $n + 2^r$ term doesn't improve.
- If we choose $r > \lg n$, then $n + 2^r$ term gets big. Example: $r = 2 \lg n \Rightarrow 2^r = 2^{2 \lg n} = (2^{\lg n})^2 = n^2$.

Radix Sort Example

- Problem: sort 1 million 64-bit numbers
 - Treat as four-digit radix 2^{16} numbers
 - $r = 16$, $d = 4$, $n = 10^6$, and $b = 64$
 - $b/r * (n + 2^r) = 4,262,144 \approx 4n$
 - We can sort in just four passes with radix sort!
- Compares well with typical $O(n \lg n)$ comparison sort
 - Requires approximately $\lg n = 20$ operations per number being sorted
 - So why would we ever use anything but radix sort?
 - Doesn't sort in place (why?)
 - Depends on implementation, e.g., quicksort uses cache better

Summary: Radix Sort

- Assumption: input has d digits ranging from 0 to k
 - Basic idea:
 - Sort elements by digit starting with least significant
 - Use a stable sort (like counting sort) for each stage
 - Each pass over n numbers with d digits takes time $O(n+k)$, so total time $O(dn+dk)$
 - When d is constant, and $k=O(n)$, takes $O(n)$ time
 - Fast, stable, and Simple to code
 - Doesn't sort in place
 - Depends on implementation, e.g., quicksort uses cache better
 - Cannot easily sort floating point numbers

Bucket Sort

- Assumes the input is generated by a random process that distributes elements uniformly over $[0, 1)$.
- ***Idea:***
 - Divide $[0, 1)$ into n equal-sized *buckets*.
 - Distribute the n input values into the buckets.
 - Sort each bucket.
 - Then go through buckets in order, listing elements in each one.

Bucket Sort (cont'd)

- Input:
 - $A[1 \dots n]$, where $0 \leq A[i] < 1$ for all i .
- Auxiliary array:
 - $B[0 \dots n - 1]$ of linked lists, each list initially empty.

Bucket sort Implementation

BUCKET-SORT(A, n)

for $i \leftarrow 1$ to n

do insert $A[i]$ into list $B[\lfloor n \cdot A[i] \rfloor]$

for $i \leftarrow 0$ to $n - 1$

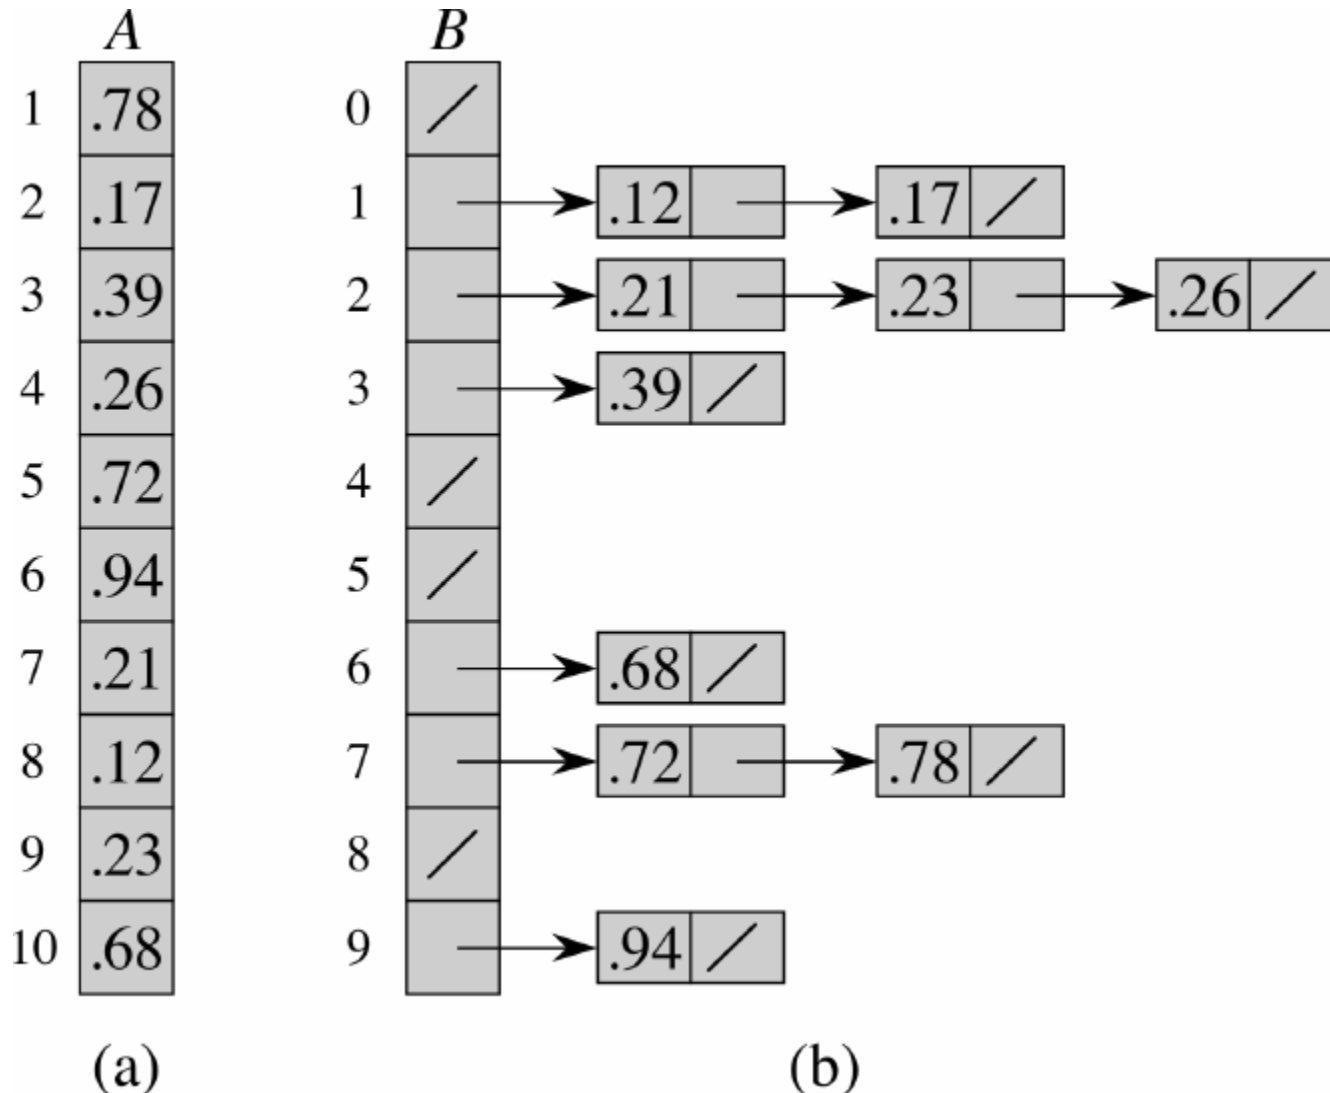
do sort list $B[i]$ with insertion sort

concatenate lists $B[0], B[1], \dots, B[n - 1]$ together in order

return the concatenated lists

Easily compute the bucket index $\lfloor n \cdot A[i] \rfloor$

Bucket sort with 10 buckets



Correctness

- Consider $A[i]$ and $A[j]$
 - Assume without loss of generality that $A[i] \leq A[j]$
 - Then, bucket index $n \cdot A[i] \leq n \cdot A[j]$
- $A[i]$ is placed into the same bucket as $A[j]$ or into a bucket with a lower index
 - If same bucket, insertion sort fixes up
 - If earlier bucket, concatenation of lists fixes up

Informal Analysis

- All lines of algorithm except insertion sorting take $\Theta(n)$ altogether
- Since the inputs are uniformly and independently distributed over $[0,1)$, we do not expect many numbers to fall into each bucket
- Intuitively, if each bucket gets a constant number of elements, it takes $O(1)$ time to sort each bucket $\Rightarrow O(n)$ sort time for all buckets.

Formal Analysis

- Define a random variable:
 n_i = the number of elements placed in bucket $B[i]$
- Because insertion sort runs in quadratic time, bucket sort time is

$$T(n) = \Theta(n) + \sum_{i=0}^{n-1} O(n_i^2) .$$

Formal Analysis (Cont'd)

Take expectations of both sides:

$$\begin{aligned} E[T(n)] &= E\left[\Theta(n) + \sum_{i=0}^{n-1} O(n_i^2)\right] \\ &= \Theta(n) + \sum_{i=0}^{n-1} E[O(n_i^2)] \quad (\text{linearity of expectation}) \\ &= \Theta(n) + \sum_{i=0}^{n-1} O(E[n_i^2]) \quad (E[aX] = aE[X]) \end{aligned}$$

n_i = the number of elements placed in bucket $B[i]$

n_i = the number of elements placed in bucket $B[i]$

Claim

$$E[n_i^2] = 2 - (1/n) \text{ for } i = 0, \dots, n-1.$$

Proof of claim

Define indicator random variables:

- $X_{ij} = I\{A[j] \text{ falls in bucket } i\}$
- $\Pr\{A[j] \text{ falls in bucket } i\} = 1/n$

- $n_i = \sum_{j=1}^n X_{ij} \quad \begin{aligned} X_{ij} &= I\{A[j] \text{ falls in bucket } i\}. \\ &= \begin{cases} 1 & \text{if } A[j] \text{ falls in bucket } i \\ 0 & \text{if } A[j] \text{ doesn't fall in bucket } i \end{cases} \end{aligned}$

The Claim

Then

$$\begin{aligned}
 E[n_i^2] &= E\left[\left(\sum_{j=1}^n X_{ij}\right)^2\right] && (x_1+x_2+x_3)(x_1+x_2+x_3) \\
 &= E\left[\sum_{j=1}^n X_{ij}^2 + 2\sum_{j=1}^{n-1}\sum_{k=j+1}^n X_{ij}X_{ik}\right] && = x_1^2 + x_1x_2 + x_1x_3 \\
 &&& + x_2^2 + x_1x_2 + x_2x_3 \\
 &&& + x_3^2 + x_1x_3 + x_2x_3 \\
 &= \sum_{j=1}^n E[X_{ij}^2] + 2\sum_{j=1}^{n-1}\sum_{k=j+1}^n E[X_{ij}X_{ik}] \quad (\text{linearity of expectation}) \\
 E[X_{ij}^2] &= 0^2 \cdot \Pr\{A[j] \text{ doesn't fall in bucket } i\} + 1^2 \cdot \Pr\{A[j] \text{ falls in bucket } i\} \\
 &= 0 \cdot \left(1 - \frac{1}{n}\right) + 1 \cdot \frac{1}{n} \\
 &= \frac{1}{n}
 \end{aligned}$$

Analysis

$E[X_{ij}X_{ik}]$ for $j \neq k$: Since $j \neq k$, X_{ij} and X_{ik} are independent random variables

$$\begin{aligned}\Rightarrow E[X_{ij}X_{ik}] &= E[X_{ij}]E[X_{ik}] \\ &= \frac{1}{n} \cdot \frac{1}{n} \\ &= \frac{1}{n^2}\end{aligned}$$

Therefore:

$$E[n_i^2] = \sum_{j=1}^n \frac{1}{n} + 2 \sum_{j=1}^{n-1} \sum_{k=j+1}^n \frac{1}{n^2}$$

Analysis (Cont'd)

$$\begin{aligned} &= n \cdot \frac{1}{n} + 2 \binom{n}{2} \frac{1}{n^2} \\ &= 1 + 2 \cdot \frac{n(n-1)}{2} \cdot \frac{1}{n^2} \\ &= 1 + \frac{n-1}{n} \\ &= 1 + 1 - \frac{1}{n} \\ &= 2 - \frac{1}{n} \end{aligned}$$

■ (claim)

Therefore:

$$\begin{aligned} \mathbb{E}[T(n)] &= \Theta(n) + \sum_{i=0}^{n-1} O(2 - 1/n) \\ &= \Theta(n) + O(n) \\ &= \Theta(n) \end{aligned}$$

Analysis conclusion

- This is a probabilistic analysis
 - We used probability to analyze an algorithm whose running time depends on the distribution of inputs.
- With bucket sort, if the input isn't drawn from a uniform distribution on $[0, 1)$, all bets are off
 - Performance-wise, but the algorithm is still correct

Bucket Sort Summary

- Assumption: input is n real #'s from $[0, 1)$
 - We can map other number into the range of $[0, 1)$
- Basic idea:
 - Create n linked lists (*buckets*) to divide interval $[0,1)$ into subintervals of size $1/n$
 - Add each input element to appropriate bucket and sort buckets with insertion sort
- Uniform input distribution $\rightarrow O(1)$ bucket size
 - Therefore the expected total time is $O(n)$

Linear Sorting Common Mistakes

- Using counting sort, when memory is limited
 - The size of $k \rightarrow$ the size of $C[0..k]$
- Using bucket sort, when the input are not uniform distributed

Linear-time Sorting Summary

- We have learned three linear-time sorting algorithms
- Their assumptions on input
 - Counting sort $\rightarrow [0..k]$
 - Radix sort $\rightarrow d$ digits
 - Bucket sort \rightarrow uniform distribution $[0, 1)$